Workshop 6 - Paging: problem, files, and tests

Team Cardinal
CSCI 3453 Spring 2017

April 22, 2017

1 Paging

1.1 Overview

Paging is a virtual memory management mechanism by which a computer stores and retrieves data from secondary storage (disk, harddrive, etc.) for use in main memory (RAM). Through this, the operating system retrieves data from secondary storage in continuous blocks of uniform size called pages. This size varies between operating systems, but within the Pintos 80x86 architecture a page is specified as 4096 bytes.

Paging is a necessary part of virtual memory since the sum of each active programs required memory exceeds the physical memory available to the system. By using secondary storage, active programs memory requirements can exceed the actual size of available physical memory.

Each of Pintos' processes have independent pages which are allocated below virtual address PHYS_BASE. The set of kernel pages, on the other hand, are located above PHYS_BASE. Kernel pages are global and fixed-size, remaining the same regardless of what thread or process is active. The kernel can access both user and kernel pages whereas a user process is restricted to user pages.

Each virtual page beginning is placed at memory addresses that are multiples of the page size. This strict placement is referred to as page alignment and allows for efficient and fast access to memory.

When a program is about to execute, the operating system has to divide its associated memory into pages, and then transfer the pages into RAM to begin execution. To the user, the entire program appears to occupy contiguous space in RAM at all times. In reality, not all pages of a program are necessarily in RAM, and the pages most likely dont occupy contiguous space. When that virtual page becomes inactive or the allocated space is needed for a different page, it is pushed to secondary storage. Poor use of locality and exhaustive use of the physical memory can result in a constant state of paging, which is more accurately called thrashing. This is the process by which the operating system must rapidly exchange pages from secondary storage to continue executing a program. Since the speed of retrieval from disk is orders of magnitude slower than RAM or even cache, the user would notice considerable declined performance, as more time is spent paging than executing.

In P3, Pintos requires that paging only be implemented when a page fault occurs and a program requires a page that isnt currently in virtual memory. Furthermore, Pintos doesnt

require that we preempt pages into cache to anticipate need. This document discusses the subproblems that must be tackled in order to implement paging within Pintos.

2 Paging Subproblems

2.1 Segments Loaded from Executables

Pintos is designed to load executable files in a lazy manner. When a process is first loaded from memory there are two ways we can go about loading all the necessary information needed to run the program. First, we could load the process and in addition, load each page that the process will need to access throughout the course of the program execution. An alternate method is load the process and only load associated pages when the process actually needs access to them. The main benefit to lazy loading is that each process experiences less latency when initially starting up since less information is being loaded from the disk. Because of this, more processes are able to be loaded at once. The cost of this is small delays in each program each time a new page is accessed.

The most notable consequence of this design is what signal the operating system has that a new page needs to be loaded into memory. As the program runs, it will encounter a page fault every time it attempts to access a page which has not yet been loaded into main memory. The Pintos kernel is then tasked with intercepting the page fault and loading the appropriate page into memory before the process can continue, evicting a page or pages as necessary to create room in main memory. A major design consideration in Project 3 will be how to handle page faults as they arise during the course of program execution but in general the following actions must be taken by the page fault handler:

1. Look up the page that faulted in the Pintos supplemental page table

    (a) Verify that the memory reference is in a valid address space
    (b) Terminate process if access to that page is invalid (read-only, kernel addr. space etc.)

2. Obtain a new frame to store the page in main memory

    (a) May lead to eviction of page currently loaded in memory

3. Read the data from disk/swap into the newly acquired frame

4. Link the new physical page to the faulty address in the page table to subsequent access can be made

Accomplishing this will require modification of the page fault handler code. The *page_fault()* function is located in *thread/exception.c*. Additionally, functions to handle adding the address to the page table can be found in *userprog/pagedir.c*.

2.2 Page Replacement - Approximating Least Recently Used (LRU)

Pintos hardware provides some assistance for implementing page replacement algorithms, through a pair of bits in the page table entry (PTE) for each page. On any read or write to a page, the CPU sets the accessed bit to 1 in the pages PTE, and on any write, the CPU sets the dirty bit to 1. The CPU never resets these bits to 0, but the OS may do so.

It is important to consider that two (or more) pages can refer to the same frame; this is known as aliasing. When an aliased frame is accessed, the accessed and dirty bits are updated in only one page table entry (the one for the page used for access). The accessed and dirty bits for the other aliases are not updated. In Pintos, every user virtual page is aliased to its kernel virtual page. A design implementation could check and update the accessed and dirty bits for both addresses, which are:

1. PTE_A Bit 5, the accessed bit.

2. PTE_D Bit 6, the dirty bit.

Refer to A.7.4.2 of the Pintos manual for more details.

The global page replacement algorithm is generally designed as follows:

Maintain a circular list of pages keeping a running count of how often a page is accesses or referenced. Utilize both the accessed and dirty bits to determine replacement policy. A pointer traverses the circular linked list looking for the accessed bit equal to 0. If the current page accessed bit is equal to 0, then check the dirty bit. If the dirty bit is also equal to 0, replace the page. If the dirty bit is 1, then reset the bit to 0 and then advance the pointer for evaluation at the next page. This policy makes use of the dirty bit as a second chance due to the expensive cost of replacing dirty pages. In threads/pte.h, it goes over in greater details convenience functions, defines, and comments on the use of the PTE. It is also important to note that the page replacement algorithm will also be used to manage the frame table.

2.3 Page Fault Synchronization

If a page fault requires I/O, it must be able to read data from the file system and/or the swap slot concurrently with currently non-faulting processes and page faults that do not require I/O. Thus, to prevent file system corruption & to be mindful of functioning processes which may require I/O at any time, page faults that need to read data from the file system must acquire the global file system lock. This requires refactoring *page_fault(struct intr_frame *f)* in *src/userprog/exception.c* to determine if the type of page fault makes use of the Supplemental Page Table implemented for this project. By default, the *page_fault()* function provides three variables that determine the cause:

1. *bool not_present* - True if the page fault was triggered by an unloaded page, false if the page fault was triggered by writing to a read-only page.

2. *bool write* - True if the process was writing when the page fault was triggered, false if the process was reading when the page fault was triggered.

3. *bool user* - True if the page fault was triggered by a user-level process, false if the page fault was triggered by a kernel-level process.

Logically speaking, if *not_present* is true and the page fault was triggered by a user level process, you can then rely on the Supplemental Page Table to determine if the file system lock should be acquired & thus synchronize the I/O-requiring page fault.

2.4   Extension of the Program Loader

The core of the program loader must be modified. The program loader is the loop found in the load_segment() function.

Each time around the loop:

- page_read_bytes receives the number of bytes to read from the executable file

- page_zero_bytes receives the number of bytes to initialize to zero, following the bytes read PGSIZE is always the sum of page_read_bytes and page_zero_bytes and is always equal to 4,096

The page handling process depends on the above variables values. Cases and corresponding page handling is as follows:

- If page_read_bytes equals PGSIZE

  - The page should be demand paged from the underlying file on its first access.

- If page_zero_bytes equals PGSIZE

  - The page does not need to be read from the disk at all (because it contains all zeroes).
  - Create a new page consisting of all zeroes at the first page fault.

- If neither page_read_bytes or page_zero_bytes equals PGSIZE

  - An initial part of the page must be read from the underlying file and the remainder must be zeroed.

3   Pintos Paging Tests

Here's a list of the tests relating to paging:
1) All the regression testing of P2 and its file system access (*src/tests/userprog*)
2) The following tests in *src/tests/vm*:

- page-linear.c

- page-merge-mm.c

- page-merge-par.c

- page-merge-seq.c

- page-merge-stk.c

- page-parallel.c

- page-shuffle.c

Open Questions

1. Should page faults that require I/O be blocked until all non-faulty running processes and non-I/O page faults finish execution, or should their precedence be determined with a round-robin scheduler? In other words, is starvation a concern with page fault synchronization?

2. Why is a "second change" LRU implementation preferential to a FIFO implementation?

3. If a page has already been loaded by one process, how can we ensure that the second process does not page fault and attempt to load that page into memory a second time?

4. Should there be a synchronization construct that locks the swap slot when in use?

Conclusion

Proper implementation and appreciation of the data structures detailed in the Resource Management Overview of the Pintos manual (section 4.1.3) requires a fundamental understanding of the purpose of (and problems posed by) paging. Since paging comprises a large portion of Pintos' third project (requiring usage of both the provided page directory and our own page loading functionality), adherence to well-documented design will streamline the project and allow teams to split up the paging subproblems amongst their members.

References

- Pintos Source Code

  - userprog/process.c
    * load_segment()
  - userprog/exception.c
    * page_fault()

- Operating Systems: Principles and Practices Second Edition by Thomas Anderson and Michael Dahlin

- Pintos Manual by Ben Pfaff